

# **BYZANTINE AGREEMENT AND PAXOS**


# Consensus or Byzantine Agreement

- Malicious Failures (byzantine failures)
- General sends a binary value to  $n-1$  participants such that:
  1. Agreement: All correct participants agree on same value
  2. Validity: If general is correct, every participant agrees on the value general sends

# PAXOS

Thanks for slides: Idit Kaider, John Ousterhout and Diego Angaro

# Paxos




- Lamport the archeologist and the “Part-time Parliament” of Paxos:
  - The Part-time Parliament, TOCS 1998
  - Paxos Made Simple, *ACM SIGACT News* 2001.
  - Paxos Made Live, PODC 2007
  - .....

# Distributed State Machine



- Fault-tolerance through replication.
  - Need to ensure that replicas remain consistent.
  - Replicas must process requests in the same order.

# Goal: Replicated Log



- Replicated log => replicated state machine
  - All servers execute same commands in same order
- Consensus module ensures proper log replication

# Correctness

- **Safety**
  - Only a value that has been proposed may be chosen.
  - Only a single value is chosen.
  - A node never learns that a value has been chosen unless it actually has been.
- **Liveness**
  - Some proposed value is eventually chosen.
  - If a value has been chosen, a node can eventually learn the value

# Paxos System Assumptions

- Paxos is an **asynchronous** consensus algorithm
  - Asynchronous networks
- Set of processes is **known a-priori**
- Failure model: **fail-stop** (not Byzantine),  
**delayed/lost messages**




# Paxos Properties

- Paxos guarantees **safety**.
  - Consensus is a **stable property**: once reached it is never violated; the agreed value is not changed.
- Paxos does **not** guaranteed **liveness**.
  - Consensus is reached if “a large enough subnetwork...is non-faulty for a long enough time.”
  - Otherwise Paxos might never terminate.


# Paxos Participants

- **Proposer**



- Suggests values for consideration by Acceptors.
- Advocates for a client.

- **Acceptor**



- Considers the values proposed by proposers.
- Renders an accept/reject decision.


- **Learner (will ignore most of the time)**

- Learns the chosen value.

- In practice, each node will usually play all three roles.

# Strawman: Single Acceptor


- Simple (incorrect) approach:  
**ONE** acceptor chooses value
- What if acceptor crashes **after** choosing?
- Solution: **quorum**
  - Multiple acceptors (3, 5, ...)
  - Value *v* is **chosen** if accepted by **majority** of acceptors



If one acceptor crashes, chosen value still available

# Majority Quorums

- Majority / “Quorum”
  - A set of acceptors consisting of more than half of all acceptors.
- Any two quorums have a nonempty intersection.
- In a system with  $2F+1$  acceptors,  $F$  acceptors can fail and we'll be OK.

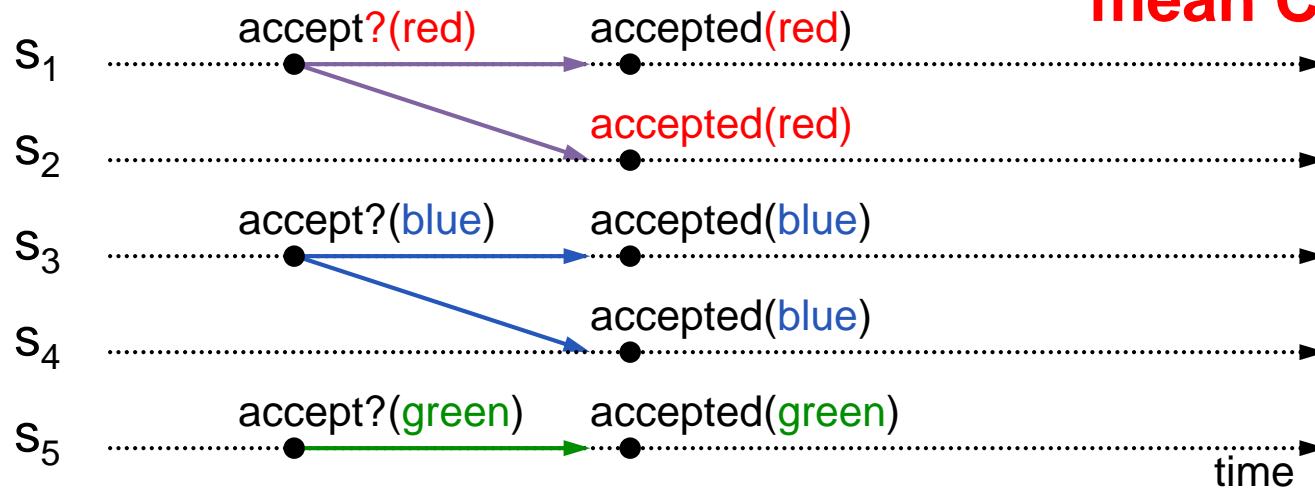


Quorums in a system with seven acceptors.

# Problem: Split Votes

- Acceptor accepts only first value it receives?
- If simultaneous proposals, no value might be chosen

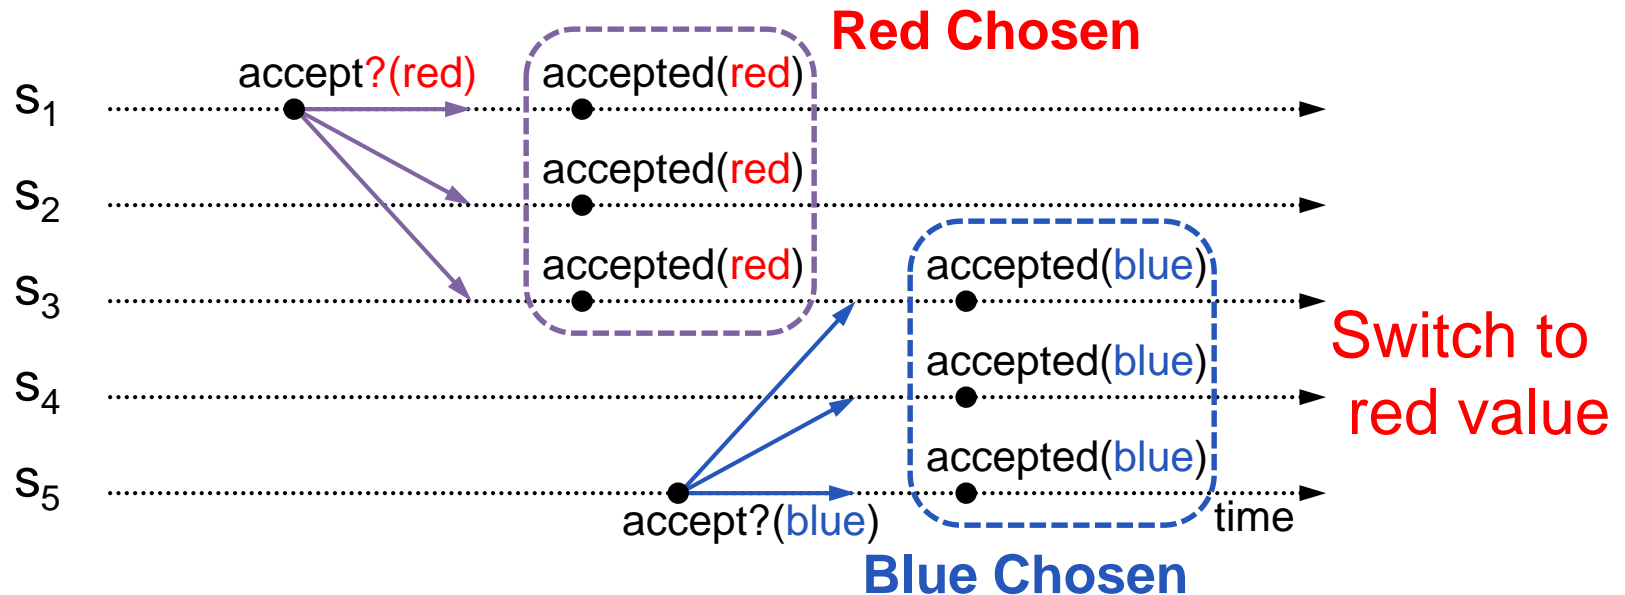
**Accepted does not mean Chosen!**



Acceptors must sometimes accept multiple (different) values (change its mind)-> multiple phases

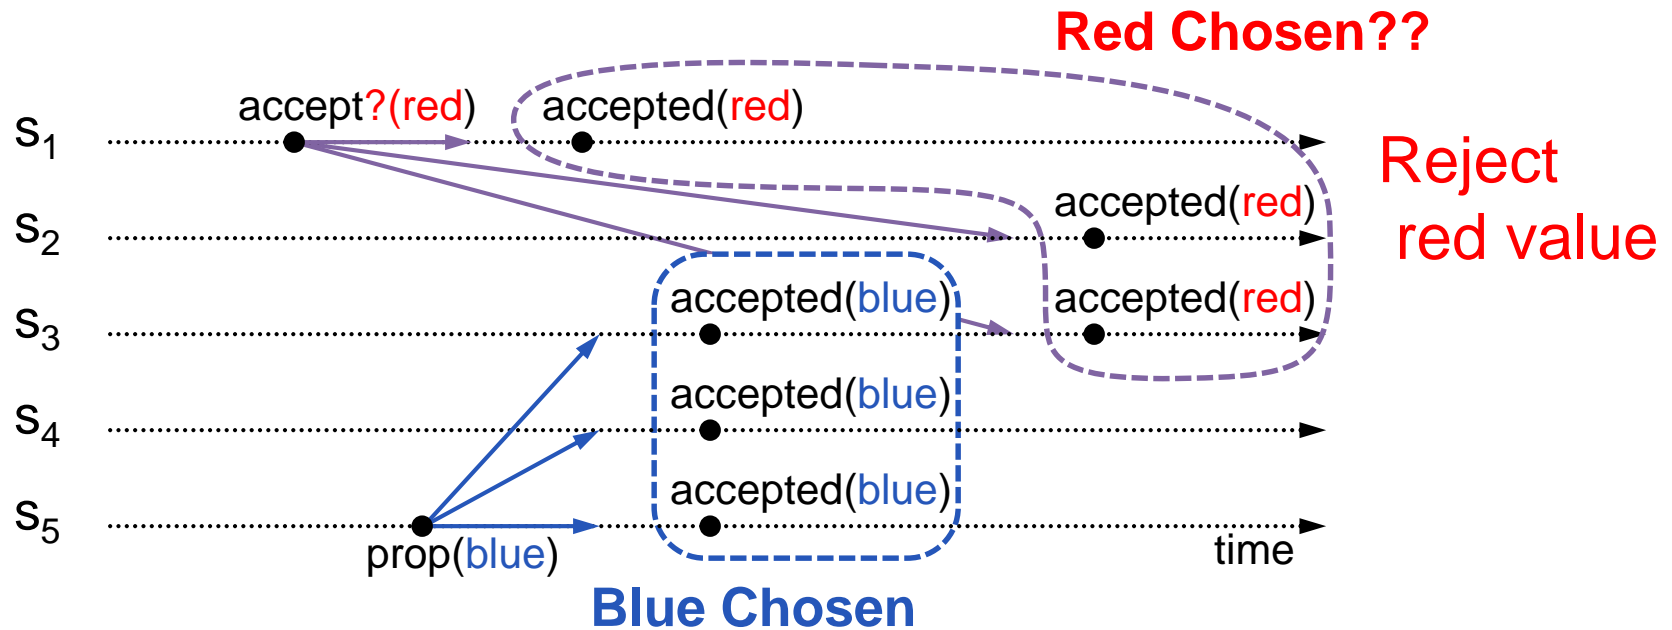
# Problem: Conflicting Choices

- What if Acceptor accepts every value it receives?



Once a value is chosen, future proposals must propose/choose the **same** value (2-phase protocol)

# Conflicting Choices, cont'd



- $s_5$  needn't propose **red** (it hasn't been chosen yet)
- $s_1$ 's proposal must be aborted ( **$s_3$  must reject it**)

**Must order proposals, reject old ones**

# Overview of the Paxos Algorithm

- **Leader based**: each client has an **estimate** of who is the **current leader**
- To order an operation, a client sends it to current leader
- The leader sequences the operation and launches an algorithm to ensure agreement



# The Consensus Algorithm Structure

- Two phases
  - Phase 1: Prepare request
  - Phase 2: Accept request
- Leader contacts a majority in each phase
- There may be multiple concurrent leaders
- A proposal consists of a unique ballot *number* and a proposed *value*.

# Ballot Numbers

- *Ballots* distinguish among values proposed by different leaders
  - **Unique**, locally monotonically increasing
  - Processes **respond only to leader with highest ballot**
- Pairs  $\langle \text{num}, \text{process id} \rangle$
- $\langle n_1, p_1 \rangle > \langle n_2, p_2 \rangle$ 
  - If  $n_1 > n_2$
  - Or  $n_1 = n_2$  and  $p_1 > p_2$
- If latest known ballot is  $\langle n, q \rangle$  then
  - $p$  chooses  $\langle n+1, p \rangle$

# The Two Phases of Paxos

- Phase 1: **prepare or Leader Election**
  - If you *want to be leader*
    - Choose **new unique ballot number**
    - Learn **outcome of all smaller ballots from majority**
- Phase 2: **accept**
  - **Leader** *proposes a value* with its ballot number
  - **Leader** gets **majority to accept** its proposal
    - A value **accepted by a majority** can be **decided**
- Phase 3: **inform**
  - Leader send final value to all

# Basic Paxos Skeleton

## Phase 1a: “Prepare”

Select proposal number  $N$  and send a ***prepare***( $N$ ) request to all acceptors.

Proposer

## Phase 1b: “Promise”

If  $N > \text{number of any previous promises or acceptances}$ ,  
- send a ***promise***( $N$ ) response

## Phase 2a: “Accept!”

If proposer received promise responses from a majority,  
- send an ***accept***( $N$ ) request to all acceptors


Acceptor

## Phase 2b: “Accepted”

If  $N \geq \text{number of any previous promise}$ ,  
accept the proposal  
- send an ***accepted*** notification to the learner


# Phase 1 – Election

- *Potential leader chooses a unique ballot id, higher than seen anything*
- *Sends to all processes*
- *Processes respond to highest ballot id*
- *If a process has in a previous round decided on a value  $v'$ , it includes value  $v'$  in its response*
- *If majority respond OK then you are the leader*




# Phase 2 – Accept

- Leader sends proposed value  $v$  to all
  - use  $v=v'$  if some process already decided and sent you its decided value  $v'$




# Phase 3 – Decision and Inform

- If leader hears a majority of OKs, it lets everyone know of the decision
- Recipients receive decision



# Which is the point of no-return?


- That is, when is consensus reached in the system?





# Which is the point of no-return?

- If/when a majority of processes hear proposed value and accept it (i.e., have respond(ed) OK!)
- Processes *may not know it yet*, but a decision has been made for the group
  - Even leader does not know it yet
- What if leader fails after that?
  - Keep having rounds until some round completes



# Paxos - Variables

BallotNum<sub>*i*</sub>, initially  $\langle 0, 0 \rangle$

Latest ballot  $p_i$  took part in (phase 1)

AcceptNum<sub>*i*</sub>, initially  $\langle 0, 0 \rangle$

Latest ballot  $p_i$  accepted a value in (phase 2)

AcceptVal<sub>*i*</sub>, initially  $\perp$

Latest accepted value (phase 2)

# Phase I: Prepare - Leader

**If I want to be leader then**

**BallotNum  $\leftarrow$   $\langle$ BallotNum.num+1, myId $\rangle$**

**send ("prepare", BallotNum) to all**

- **Goal:** contact other processes, **ask them to join this ballot**, and get information about possible past decisions

# Phase I: Prepare - Cohort

- Upon receive (“prepare”,  $bal$ ) from  $i$   
if  $bal \geq BallotNum$  then

$BallotNum \leftarrow bal$

send (“ack”,  $bal$ ,  $AcceptNum$ ,  $AcceptVal$ ) to  $i$

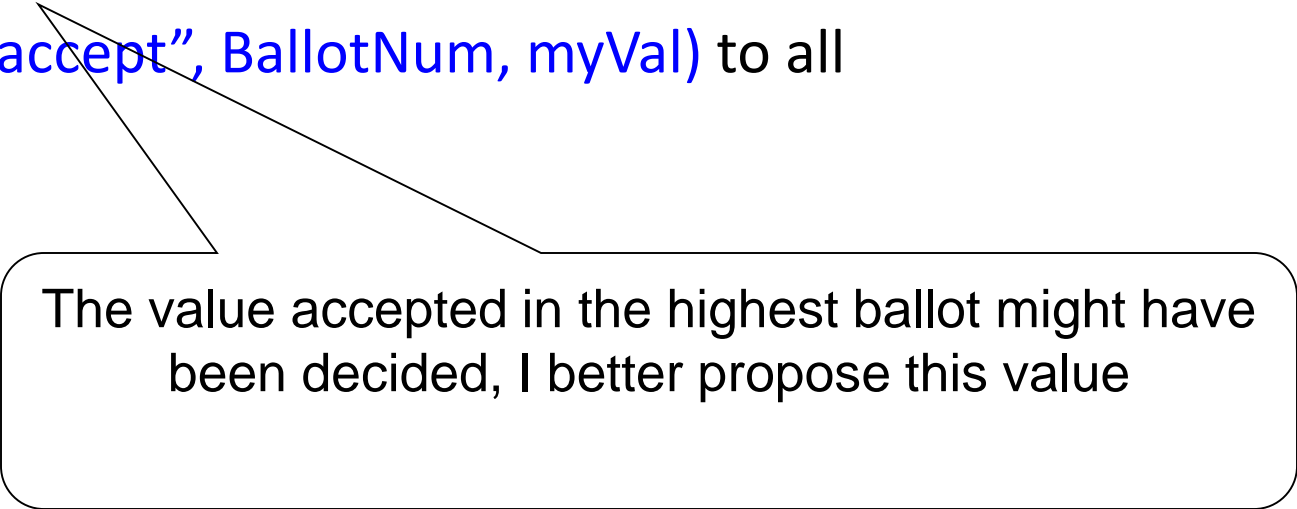
This is a higher ballot than my current, I better join it

This is a promise not to accept ballots smaller than  $bal$  in the future

Tell the leader about my latest accepted value and what ballot it was accepted in

# Phase II: Accept - Leader

Upon receive ("ack", BallotNum, b, val) from *majority*  
**if** all vals =  $\perp$  **then** myVal = initial value  
**else** myVal = received val with highest b  
send ("accept", BallotNum, myVal) to all



The value accepted in the highest ballot might have been decided, I better propose this value

# Phase II: Accept - Cohort

Upon receive (“accept”,  $b$ ,  $v$ )

**if**  $b \geq \text{BallotNum}$  **then**

$\text{AcceptNum} \leftarrow b$ ;  $\text{AcceptVal} \leftarrow v$     */\* accept proposal \*/*

send (“accept”,  $b$ ,  $v$ ) to leader

This is not from an  
old ballot

# Phase III: Decide and Inform

Upon receive (“accept”,  $b$ ,  $v$ ) from *majority*

decide  $v$

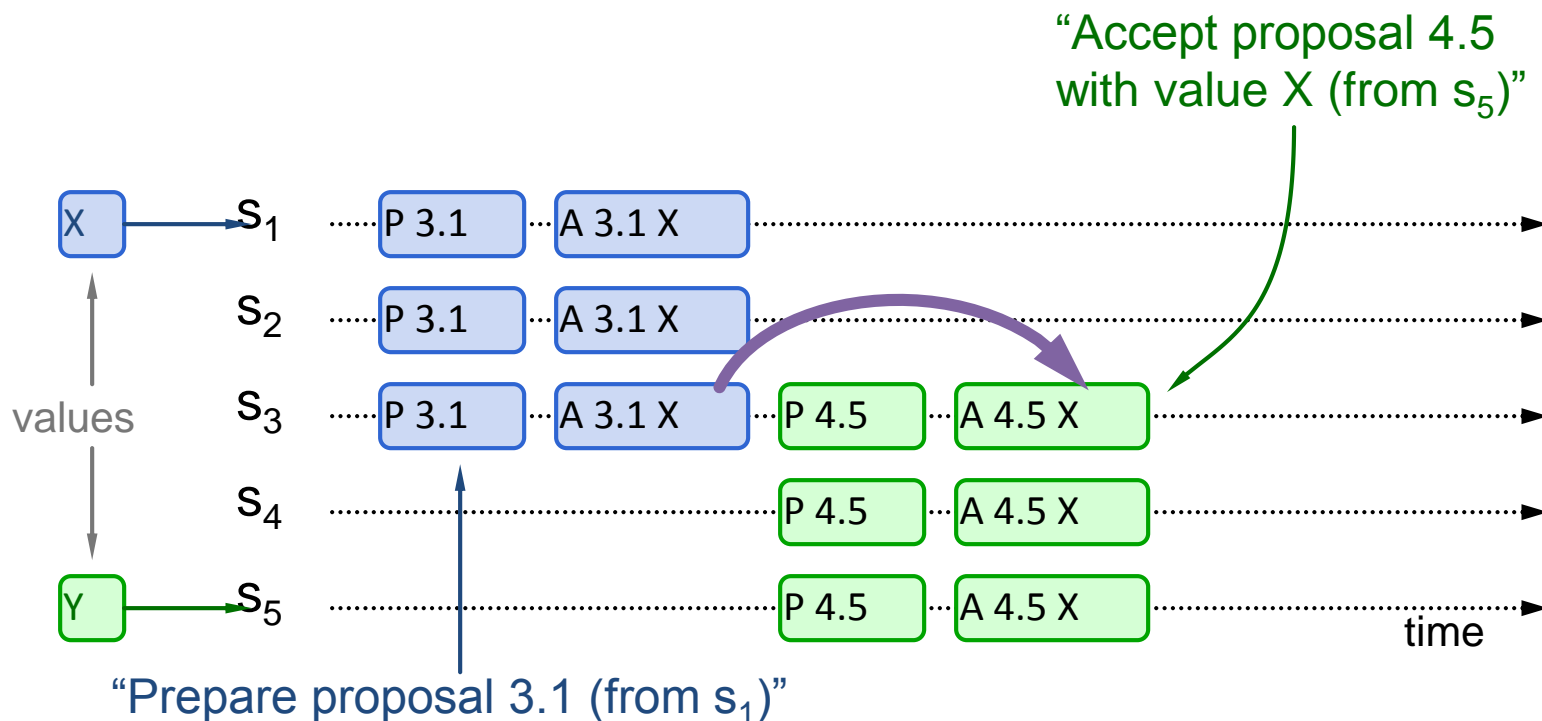
send (“accept”,  $b$ ,  $v$ ) **to all**)

# Basic Paxos Examples

Three possibilities when later proposal prepares:

## 1. Previous value already chosen:

- New proposer will find it and use it

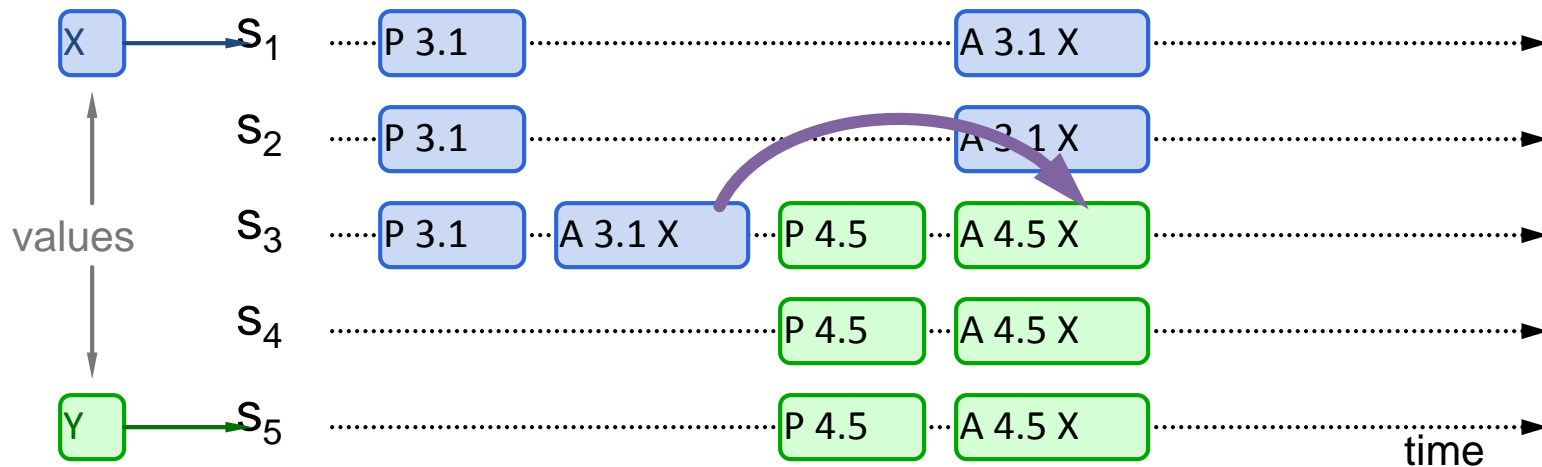




# Basic Paxos Examples, cont'd

## 2. Previous value not chosen, but new proposer sees it:

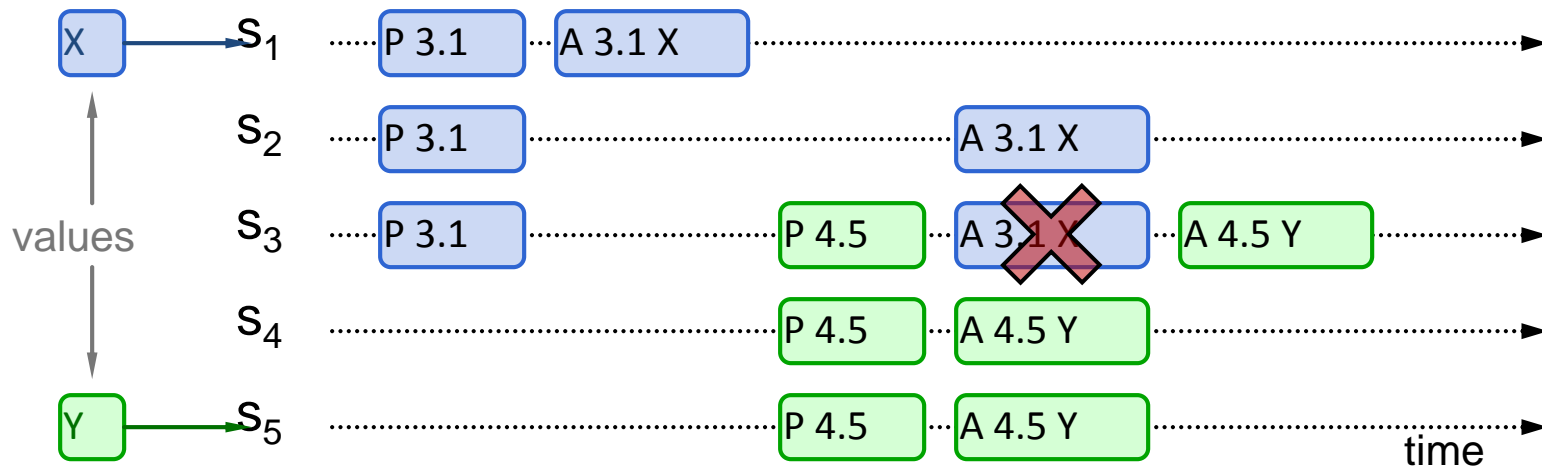
- New proposer will use existing value
- Both proposers can succeed



# Basic Paxos Examples, cont'd


## 3. Previous value not chosen, new proposer doesn't see it:

- New proposer chooses its own value
- Older proposal blocked



# Safety

- If some round has a majority hearing proposed value  $v'$  and accepting it (middle of Phase 2), then subsequently each round either:
  - 1) the round chooses  $v'$  as decision or
  - 2) the round fails
- Proof:
  - Potential leader waits for majority of OKs in Phase 1
  - At least one will contain  $v'$  (because two majorities intersect)
  - It will choose to send out  $v'$  in Phase 2
- Success requires majority, and any two majorities intersect

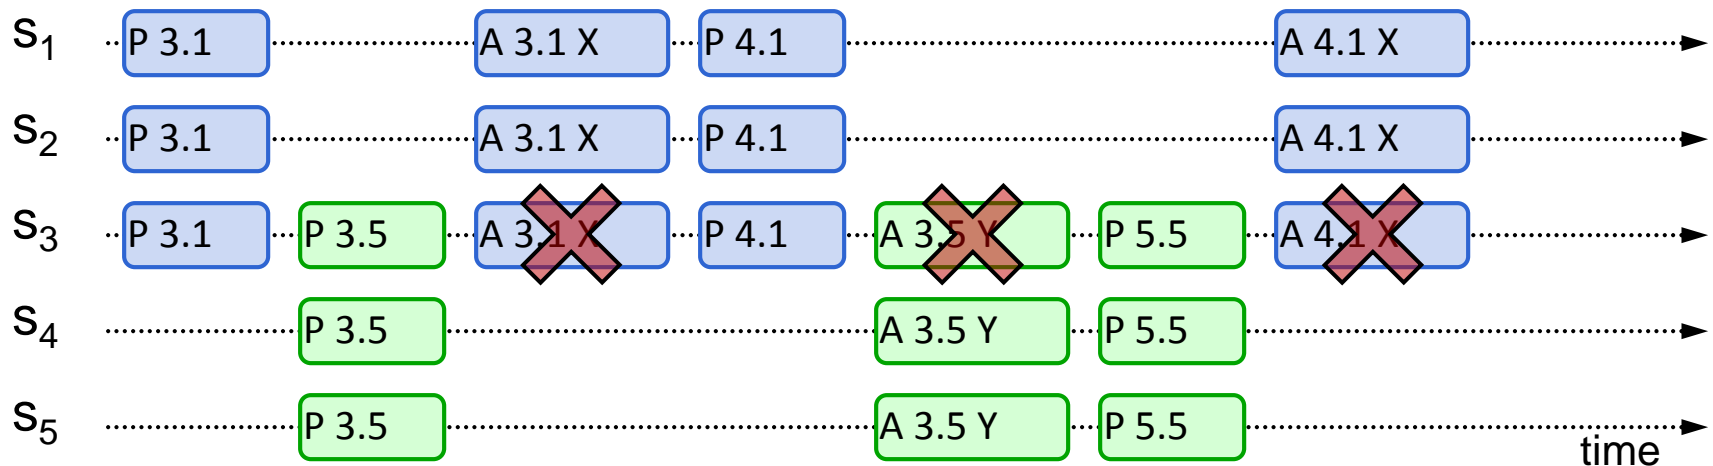


# What could go wrong?

- Process fails
  - Majority does not include it
  - When process restarts, it uses log to retrieve past decisions and past-seen ballot ids.
- Leader fails
  - Start another round
- Note that anyone can start a round any time
- Protocol may never end!
  - Impossibility result not violated
  - If things go well sometime in the future, consensus reached

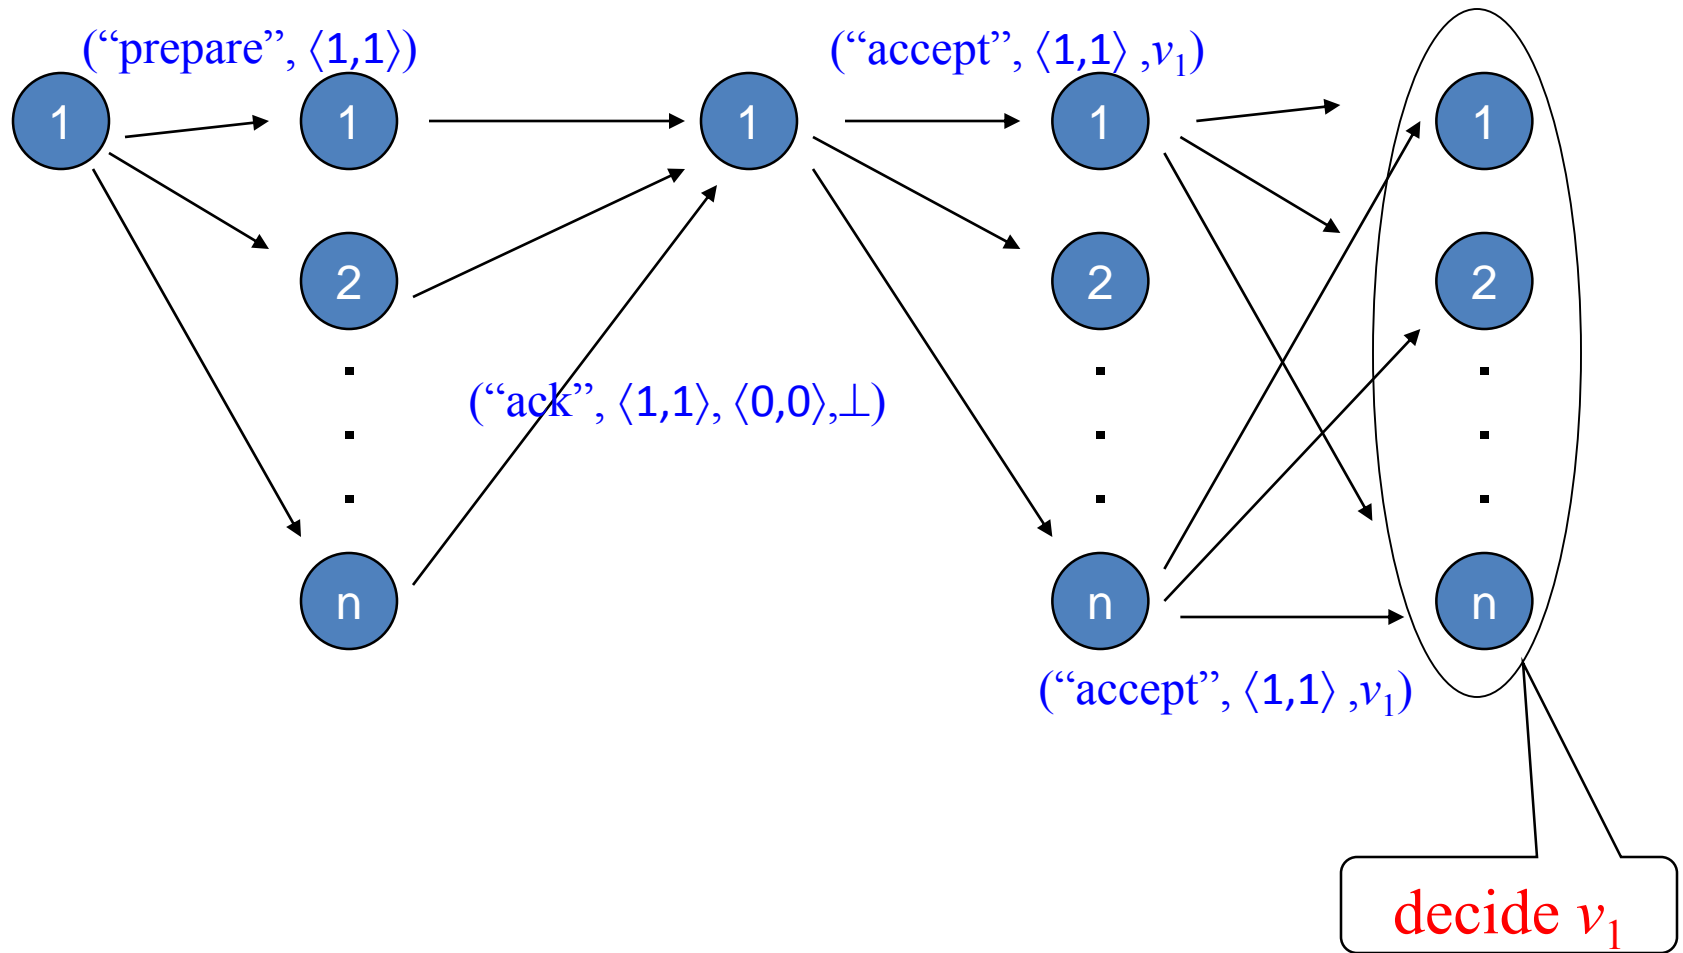
# Liveness

- Competing proposers can **livelock**:

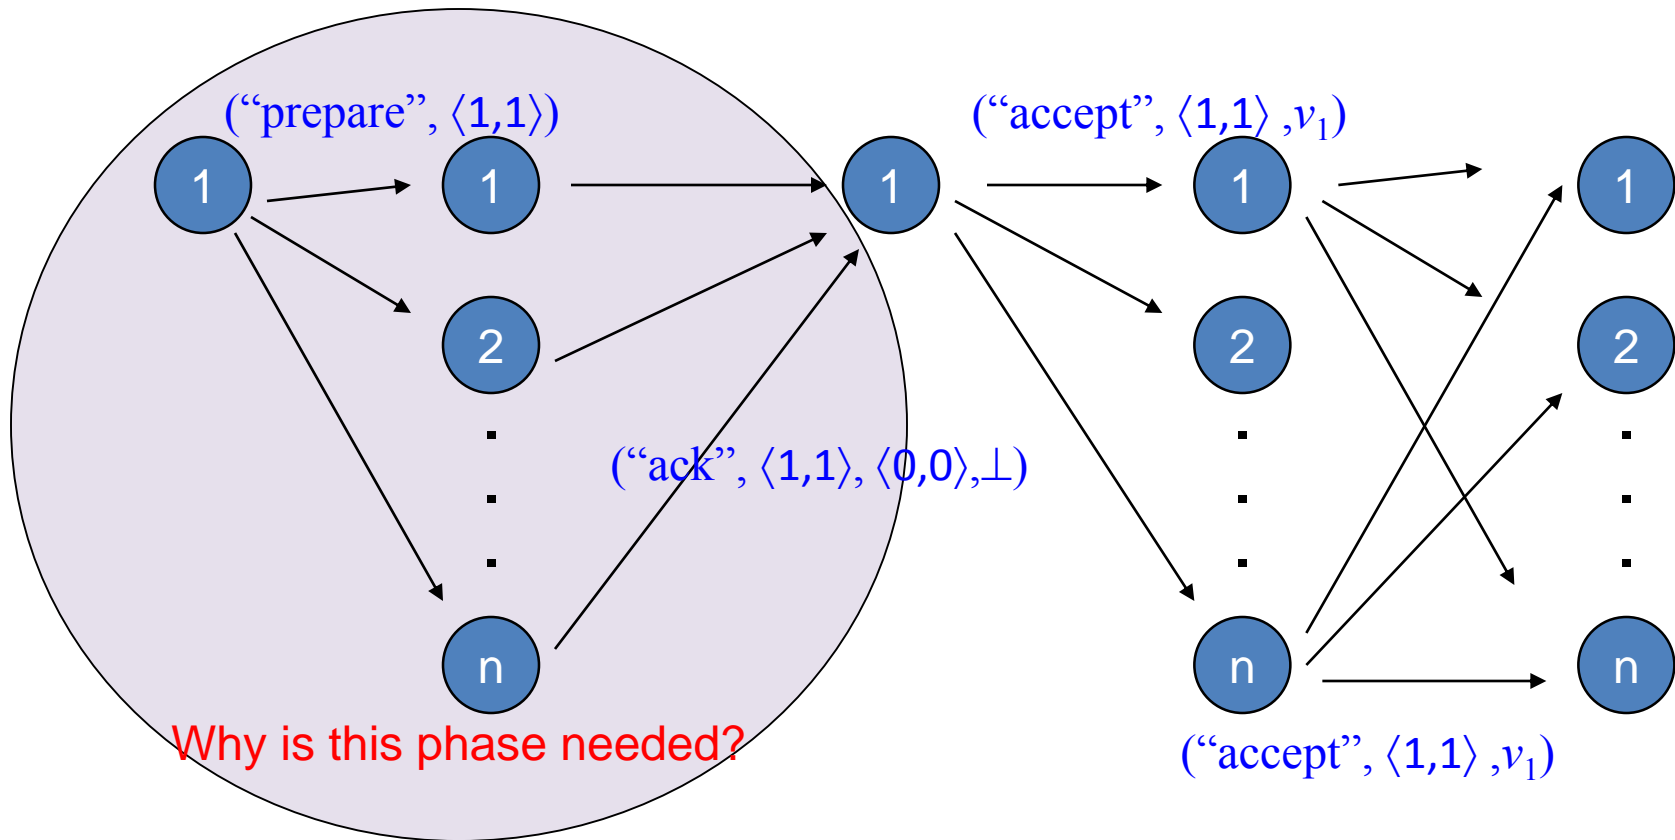


- One solution:** randomized delay before restarting
  - Give other proposers a chance to finish choosing

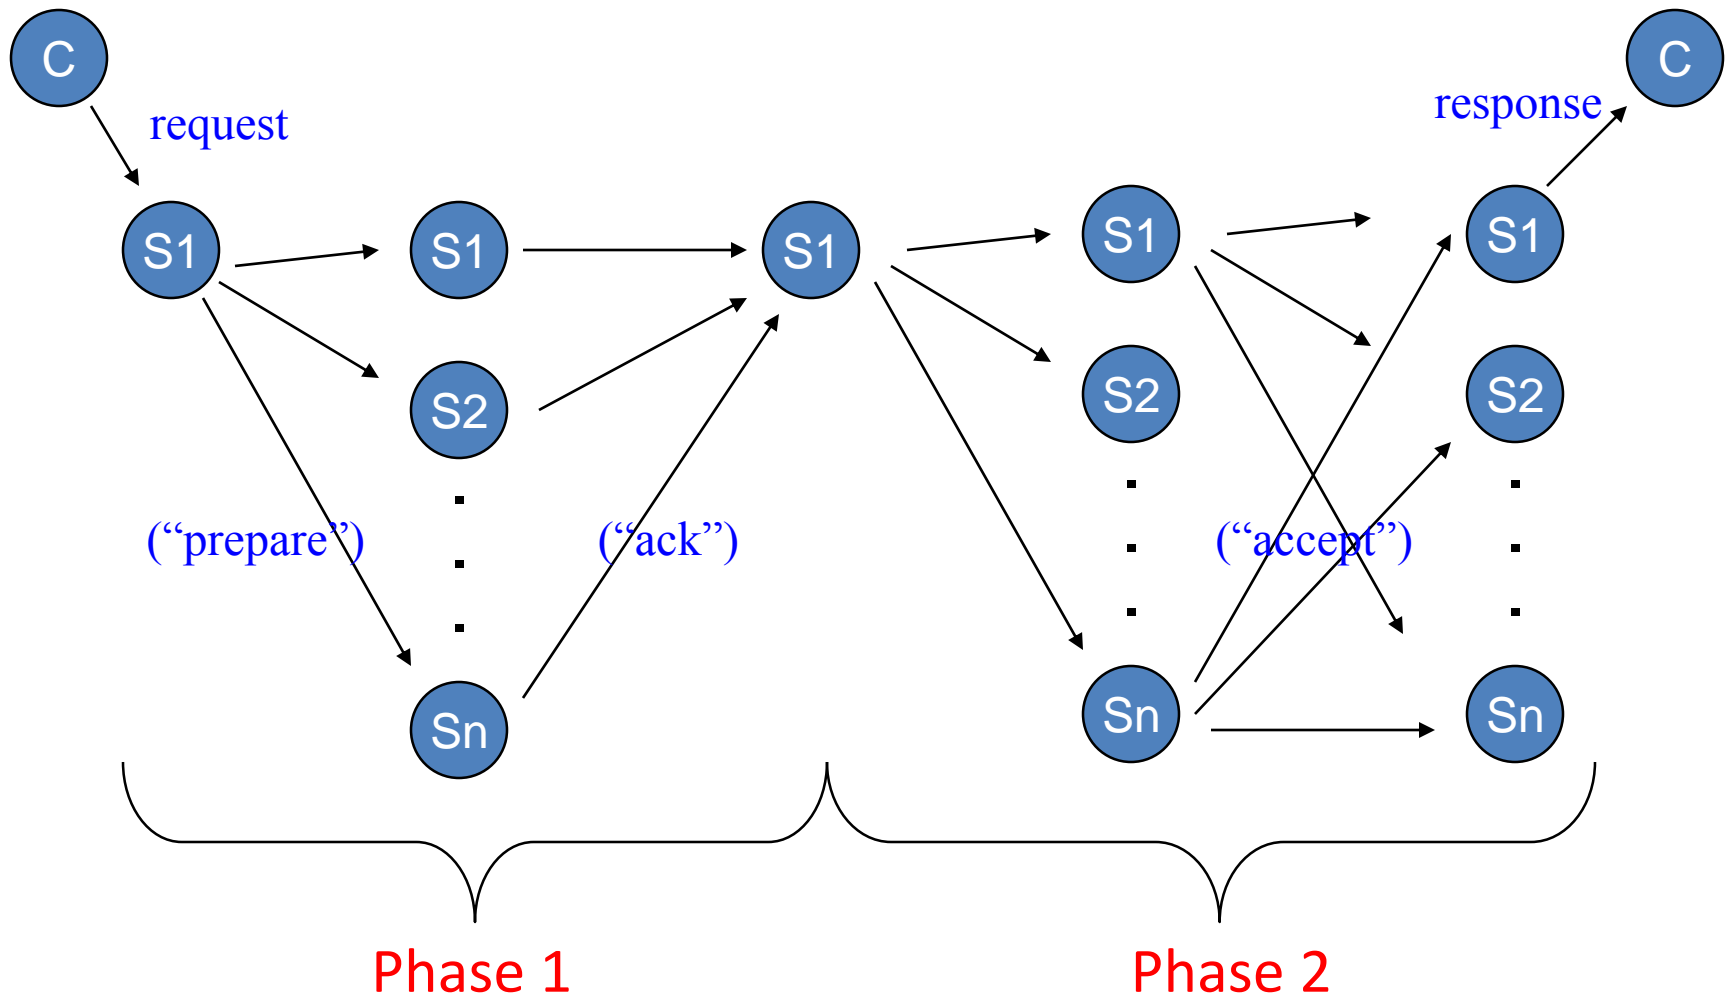
# In Failure-Free Execution



# Performance?



# Failure-Free Execution

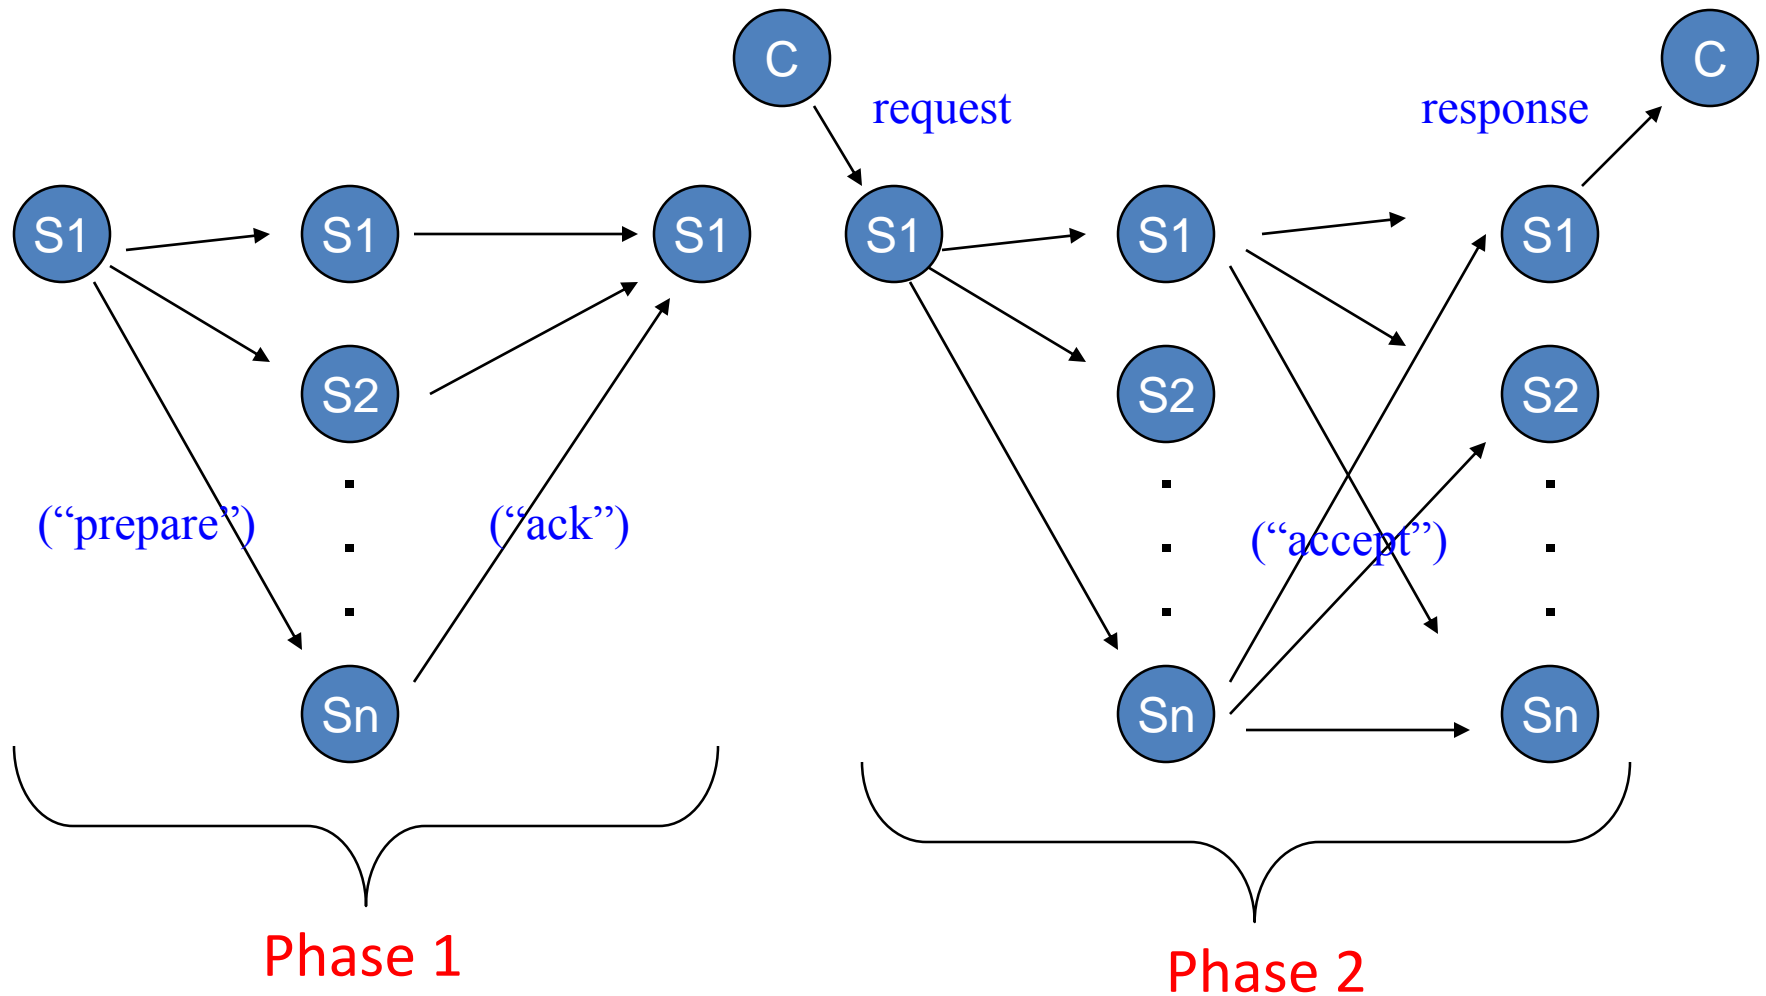




# Observation

- In Phase 1, no consensus values are sent:
  - Leader chooses largest unique ballot number
  - Gets a majority to “vote” for this ballot number
  - Learns the outcome of all smaller ballots
- In Phase 2, leader proposes its own initial value or latest value it learned in Phase 1

# Failure free execution



# Optimization: Multi-Paxos

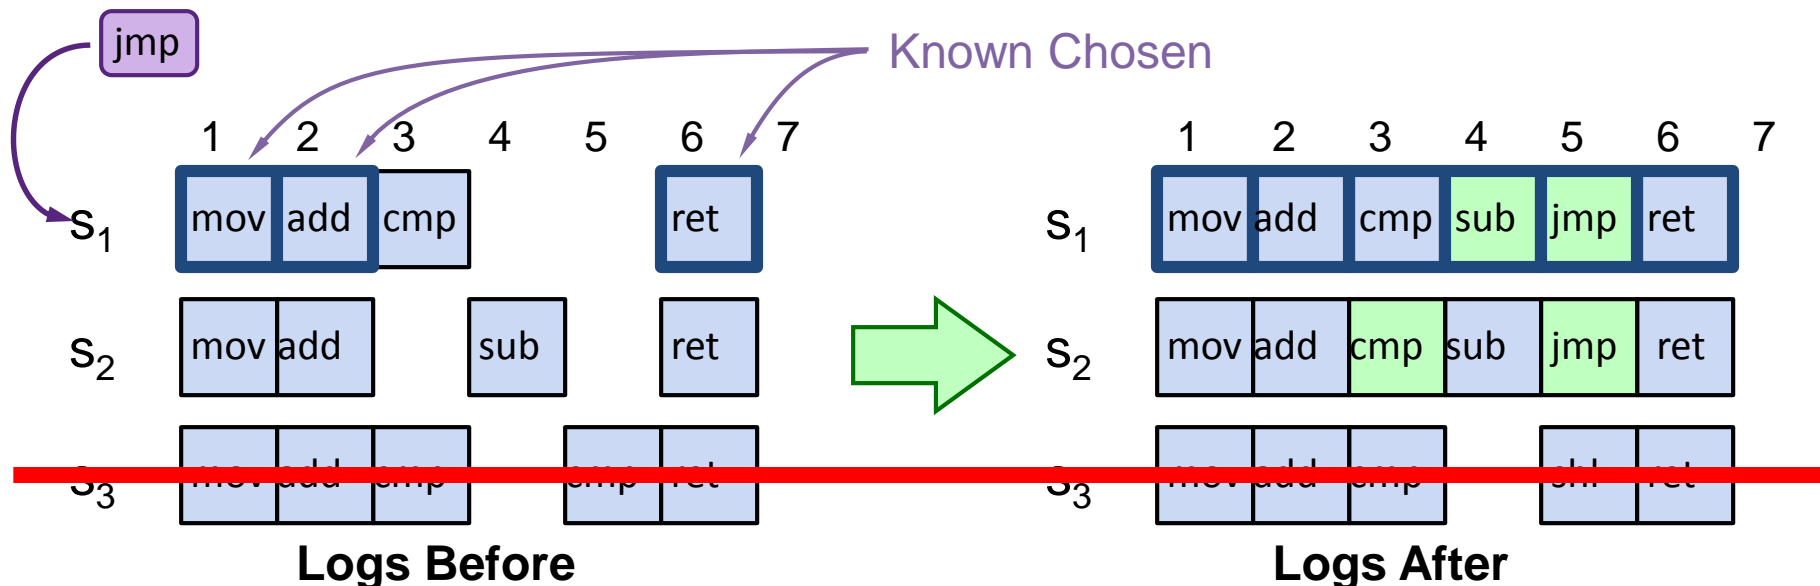
- Run **Phase 1** only when **the leader changes**
  - Phase 1 is called “**view change**” or “**leader election**”
  - Phase 2 is the “**normal mode**”
- A leader is for The Log.
- Each message includes **BallotNum** (from the last Phase 1) and **ReqNum**
- Respond only to messages with the “right” BallotNum

# Multi-Paxos

- Leader can handle multiple client requests concurrently:
  - Select different log entries for each
- Must apply commands to state machine in log order


# Selecting Log Entries

- When leader receives a request arrives from a client:
  - Find first log entry not known to be chosen
  - Run Paxos to propose client's command for this index
  - Prepare returns acceptedValue?
    - Yes: finish choosing acceptedValue, start again
    - No: choose client's command



# Configuration Changes

- Safety requirement:
  - During configuration changes, it must not be possible for different majorities to choose different values for the same log entry:




**BAD!**

# Configuration Changes, cont'd

- Use the log to manage configuration changes:
  - Configuration is stored as a log entry
  - Replicated just like any other log entry
  - Configuration for choosing entry  $i$  determined by entry  $i-\alpha$ .

Suppose  $\alpha = 3$ :



- Notes:
  - $\alpha$  determines concurrency:
    - $\alpha$  concurrent log entries.
    - But can't choose entry  $i+\alpha$  until entry  $i$  chosen
  - Issue no-op commands if needed to complete change quickly



*That's all Folks!*